

# PyGeNN: A Python library for GPU-enhanced neural networks

James C Knight<sup>1,\*</sup>, Anton Komissarov, Thomas Nowotny<sup>1</sup>

<sup>1</sup>Centre for Computational Neuroscience and Robotics, School of Engineering and Informatics, University of Sussex, Brighton, United Kingdom

Correspondence\*:

James C Knight

J.C.Knight@sussex.ac.uk

## 2 ABSTRACT

3 For full guidelines regarding your manuscript please refer to Author Guidelines.

4 As a primary goal, the abstract should render the general significance and conceptual advance  
5 of the work clearly accessible to a broad readership. References should not be cited in the  
6 abstract. Leave the Abstract empty if your article does not require one, please see Summary  
7 Table for details according to article type.

8 **Keywords:** GPU, high-performance computing, parallel computing, benchmarking, computational neuroscience, spiking neural  
9 networks, Python

## 1 INTRODUCTION

10 A wide range of spiking neural network (SNN) simulators are available, each with their own niche.  
11 NEST (Gewaltig and Diesmann, 2007) is widely used for large-scale point neuron simulations on distributed  
12 computing systems; NEURON (Carnevale and Hines, 2006) and Arbor (Akar et al., 2019) specialise in the  
13 simulation of complex multi-compartmental models; and CARLsim (Chou et al., 2018) and GeNN (Yavuz  
14 et al., 2016) use Graphics Processing Units (GPUs) to accelerate point neuron models. For performance  
15 reasons, many of these simulators are written in C++ and, especially amongst the older simulators, users  
16 describe their models either using a Domain-Specific Language (DSL) or directly in C++. For programming  
17 language purists, a DSL may be an elegant way of describing an SNN network model and, for simulator  
18 developers, not having to add bindings to another language is convenient. However, both choices act as  
19 a barrier to potential users. Therefore, with both the computational neuroscience and machine learning  
20 communities gradually coalescing towards a Python-based ecosystem with a wealth of mature libraries  
21 for scientific computing (Hunter, 2007; Van Der Walt et al., 2011; Millman and Aivazis, 2011), exposing  
22 spiking neural network simulators to Python seems a pragmatic choice. NEST (Eppler et al., 2009),  
23 Neuron (Hines et al., 2009) and CARLsim (Balaji et al., 2020) have all taken this route and now offer a  
24 Python interface. Furthermore, newer simulators such as Arbor and Brian2 (Stimberg et al., 2019) have  
25 been designed from the ground up with a Python interface.

26 While we have recently demonstrated some very competitive performance results (Knight and Nowotny,  
27 2018, 2020) using our GeNN simulator (Yavuz2016), it has not been usable directly from Python. GeNN  
28 can already be used as a backend for the Python-based Brian2 simulator (Stimberg et al., 2019) but, while  
29 Brian2GeNN (Stimberg et al., 2020) allows Brian2 users to harness the performance benefits GeNN

provides, it is not possible to expose all of GeNN’s unique features to Python through the Brian2 API. Specifically, GeNN not only allows users to easily define their own neuron and synapse models but, also ‘snippets’ for offloading the potentially costly initialisation of model parameters and connectivity onto the GPU. Additionally, GeNN provides a lot of freedom for users to integrate their own code into the simulation loop. In this paper we describe the implementation of PyGeNN – a Python package which aims to expose the full range of GeNN functionality with minimal performance overheads. We then demonstrate this performance using two large-scale models from the literature.

## 2 MATERIALS AND METHODS

### 2.1 GeNN

GeNN (Yavuz et al., 2016) is a library for generating CUDA code for the simulation of spiking neural network models. GeNN handles much of the complexity of using CUDA directly as well as automatically performing device-specific optimizations so as to maximize performance.

GeNN consists of a main library – implementing the API used to define models as well as the generic parts of the code generator – and an additional library for each backend (currently there is a reference C++ backend for generating CPU code and a CUDA backend. An OpenCL backend is under development). Users describe their model by implementing a `modelDefinition` function within a C++ file. For example, a model consisting of 4 Izhikevich neurons with heterogeneous parameters, driven by a constant input current might be defined as follows:

```

47 void modelDefinition(ModelSpec &model)
48 {
49     model.setDT(0.1);
50     model.setName("izhikevich");
51
52     NeuronModels::IzhikevichVariable::VarValues popInit(
53         -65.0, -20.0, uninitialisedVar(), uninitialisedVar(),
54         uninitialisedVar(), uninitialisedVar());
55
56     model.addNeuronPopulation<NeuronModels::IzhikevichVariable>(
57         "Pop", 4, {}, popInit);
58
59     CurrentSourceModels::DC::ParamValues csParams(10.0);
60
61     model.addCurrentSource<CurrentSourceModels::DC>(
62         "CS", "Neurons", csParams, {});
63 }
```

The `genn-buildmodel` command line tool is then used to compile this file; link it against the main GeNN library and the desired backend library; and finally run the resultant executable to generate the source code required to build a simulation dynamic library (a .dll file on Windows or a .so file on Linux and Mac). This dynamic library can then either be statically linked against a simulation loop provided by the user or dynamically loaded by the users simulation code. To demonstrate this latter approach, this example uses the `SharedLibraryModel` helper class supplied with GeNN to dynamically load the previously defined model, initialise the heterogeneous neuron parameters and print each neuron’s membrane voltage every timestep:

```

72 #include "sharedLibraryModel.h"
73
74 int main()
75 {
76     SharedLibraryModel<float> model("./", "izhikevich");
77     model.allocateMem();
78     model.initialize();
79     float *aPop = model.getScalar<float>("a");
80     float *bPop = model.getScalar<float>("b");
81     float *cPop = model.getScalar<float>("c");
82     float *dPop = model.getScalar<float>("d");
83     aPop[0] = 0.02; bPop[0] = 0.2; cPop[0] = -65.0; dPop[0] = 8.0; // RS
84     aPop[1] = 0.1; bPop[1] = 0.2; cPop[1] = -65.0; dPop[1] = 2.0; // FS
85     aPop[2] = 0.02; bPop[2] = 0.2; cPop[2] = -50.0; dPop[2] = 2.0; // CH
86     aPop[3] = 0.02; bPop[3] = 0.2; cPop[3] = -55.0; dPop[3] = 4.0; // IB
87     model.initializeSparse();
88
89     float *vPop = model.getScalar<float>("VPop");
90     while(model.getTime() < 200.0f) {
91         model.stepTime();
92         model.pullVarFromDevice("Pop", "V");
93         printf("%f, %f, %f, %f, %f\n", t, VPop[0], VPop[1], VPop[2], VPop[3]);
94     }
95     return EXIT_SUCCESS;
96 }

```

## 97 2.2 SWIG

98 In order to use GeNN from Python, both the model creation API and the `SharedLibraryModel`  
99 functionality need to be ‘wrapped’ so they can be called from Python. While this is possible using  
100 the API built into Python itself, a wrapper function would need to be manually implemented for each  
101 GeNN function to be exposed which would result in a lot of maintenance overhead. Instead, we chose  
102 to use SWIG (Beazley, 1996) to automatically generate wrapper functions and classes. SWIG generates  
103 Python modules based on special interface files which can directly include C++ code as well as special  
104 ‘directives’ which directly control SWIG:

```

105 %module(package="package") package
106 %include "test.h"

```

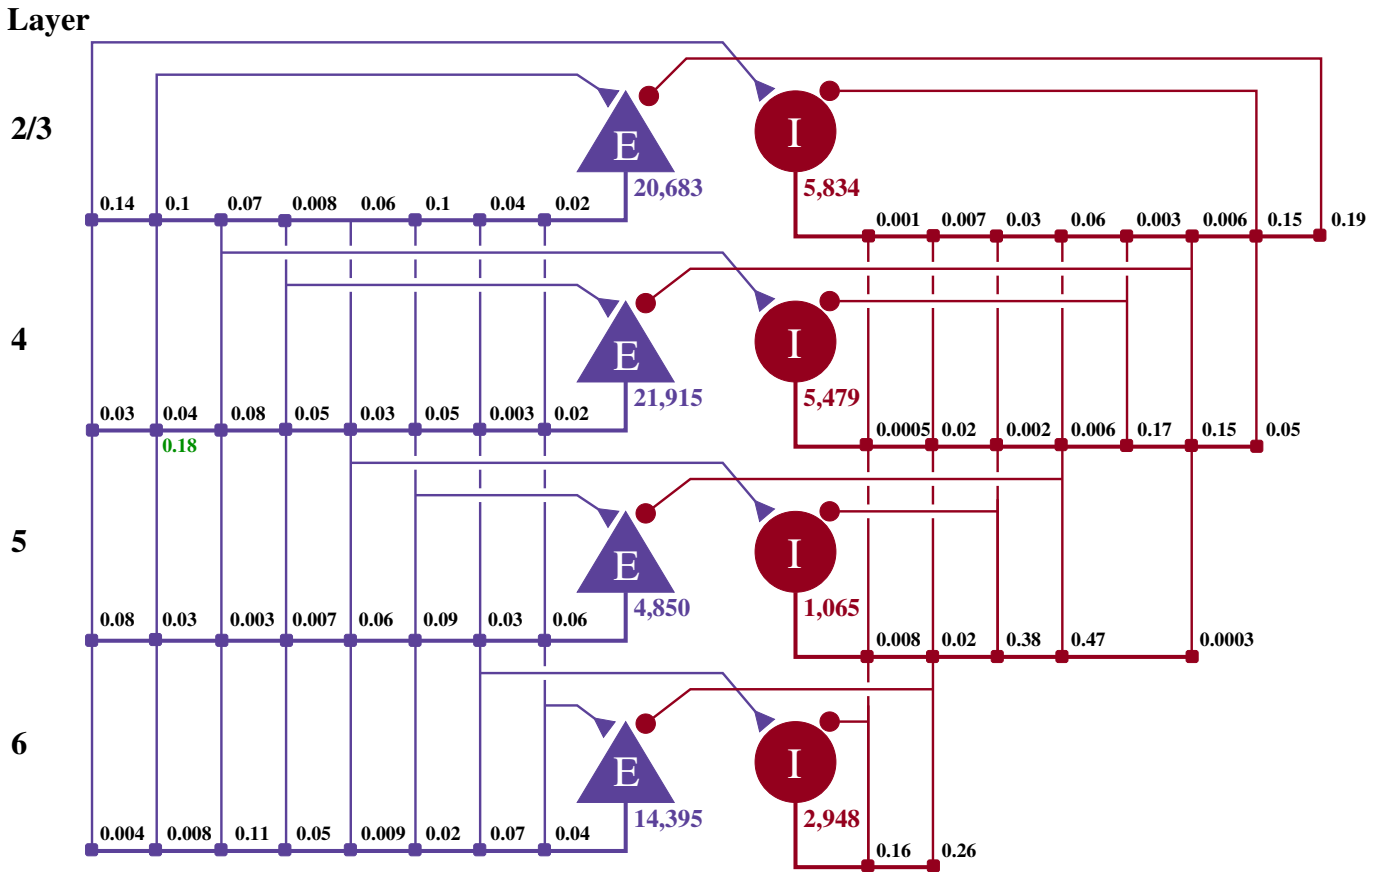
107 where the `%module` directive sets the name of the generated module and the package it will be located in  
108 and the `%include` directive parses and automatically generates wrapper functions for a C++ header file.  
109 We use SWIG in this manner to wrap both the model building and `SharedLibraryModel` APIs described  
110 in section 2.1. However, key parts of GeNN’s API such as the `ModelSpec::addNeuronPopulation` method  
111 employed in section 2.1, rely on C++ templates which are not directly translatable to Python. Instead, valid  
112 template instantiations need to be given a unique name in Python using the `%template` SWIG directive:

```

113 %template(addNeuronPopulationLIF) ModelSpec::addNeuronPopulation<NeuronModels::LIF>;

```

114 Having to manually add these directives whenever a model is added to GeNN would be exactly the sort of  
115 maintenance overhead we were trying to avoid by using SWIG. Instead, when building the Python wrapper,



**Figure 1.** Illustration of the microcircuit model. Blue triangles represent excitatory populations, red circles represent inhibitory populations and the numbers beneath each symbol shows the number of neurons in each population. Connection probabilities are shown in small bold numbers at the appropriate point in the connection matrix. All excitatory synaptic weights are normally distributed with a mean of 0.0878 nA (unless otherwise indicated in green) and a standard deviation of 0.008 78 nA. All inhibitory synaptic weights are normally distributed with a mean of 0.3512 nA and a standard deviation of 0.035 12 nA.

116 we search the GeNN header files for the macros used to declare models in C++ and automatically generate  
 117 SWIG %template directives.

118 As previously discussed, a key feature of GeNN is the ease with which it allows users to define their  
 119 own neuron and synapse models as well as ‘snippets’ defining how variables and connectivity should be  
 120 initialised. Beneath the syntactic sugar described in our previous work (Knight and Nowotny, 2018), new  
 121 models can be defined in C++ by defining a new class derived from, for example, the `NeuronModels::Base`  
 122 class. The ability to extend this system to Python was a key requirement of PyGeNN and, by using SWIG  
 123 ‘directors’, C++ classes can be made inheritable from Python using a single SWIG directive:

124 %feature("director") NeuronModels::Base;

## 125 2.3 PyGeNN

126 While GeNN *could* be used from Python via the wrapper generated using the techniques described in the  
 127 previous section, the resultant code would be unpleasant to use directly. For example, rather than being  
 128 able to specify neuron parameters using a native Python data structure such as a list or dictionary, you  
 129 have to use a wrapped type such as `DoubleVector([0.25, 10.0, 0.0, 0.0, 20.0, 2.0, 0.5])`. To provide a more

user-friendly and pythonic interface, we have built PyGeNN on top of the wrapper generated by SWIG. PyGeNN combines the separate model building and simulation stages of building a GeNN model in C++ into a single API, likely to be more familiar to users of existing Python-based model description languages such as PyNEST (Eppler et al., 2009) or PyNN (Davison et al., 2008). This allows PyGeNN to provide a unified dictionary-based API for initialising homogeneous and heterogeneous parameters as shown in this re-implementation of the previous example:

```


136 from pygenn import genn_wrapper, genn_model
137
138 model = genn_model.GeNNModel("float", "izhikevich")
139 model.dT = 0.1
140
141 izk_init = {"V": -65.0,
142            "U": -20.0,
143            "a": [0.02, 0.1, 0.02, 0.02],
144            "b": [0.2, 0.2, 0.2, 0.2],
145            "c": [-65.0, -65.0, -50.0, -55.0],
146            "d": [8.0, 2.0, 2.0, 4.0]}
147
148 pop = model.add_neuron_population("Pop", 4, "IzhikevichVariable", {}, izk_init)
149 model.add_current_source("CS", "DC", "Neurons", {"amp": 10.0}, {})
150
151 model.build()
152 model.load()
153
154 v = pop.vars["V"].view
155 while model.t < 200.0:
156     model.step_time()
157     model.pull_state_from_device("Neurons")
158     print("%t, %f, %f, %f, %f" % (model.t, v[0], v[1], v[2], v[3]))

```

While the PyGeNN API is more pythonic and, hopefully, more user-friendly than the C++ interface, it still provides users with the same low-level control over the simulation. Furthermore, by using SWIG's numpy (Van Der Walt et al., 2011) interface, the host memory allocated by GeNN can be accessed directly from Python using the `pop.vars["V"].view` syntax meaning that no potentially additional expensive copying of data is required. **(TODO: DEFINING NEW NEURON MODELS)**


## 2.4 Spike recording system

Internally, GeNN stores the spikes emitted by a neuron population during one simulation timestep in an array containing the indices of the neurons that spiked alongside a counter of how many spikes have been emitted. Previously, recording spikes in GeNN was very similar to the recording of voltages shown in the previous example code – the array of neuron indices was simply copied from the GPU to the CPU every timestep. However, especially when simulating models with a small simulation timestep, such frequent synchronization between the CPU and GPU is costly – especially if a higher-level language such as Python is involved. Furthermore, biological neurons typically spike at a low rate (in the cortex, the average firing rate is only around 3 Hz (Buzsáki and Mizuseki, 2014)) meaning that the amount of spike data transferred every timestep is typically very small. To address both of these sources of inefficiency, we have added a new data structure to GeNN which stores spike data for many timesteps on device. To reduce the memory



**Figure 2.** Illustration of the balanced random network model. The blue triangle represents the excitatory population, the red circle represents the inhibitory population, and the numbers beneath each symbol show the number of neurons in each population. Connection probabilities are shown in small bold numbers at the appropriate point in the connection matrix. All excitatory synaptic weights are initialised to 0.045 61 nA and all inhibitory synaptic weights are initialised to 0.228 05 nA.

175 required for this data structure and to make its size independent of neural activity, the spikes emitted by a  
 176 population of  $N$  neurons in a single simulation timestep are stored in a  $N$ bit bitfield where a ‘1’ represents  
 177 a spike and a ‘0’ the absence of one. Spiking data over multiple timesteps is then represented by bitfields  
 178 stored in a circular buffer. Using this approach, even the spiking output of relatively large models, running  
 179 for many timesteps can be stored in a small amount of memory. For example, the spiking output of a model  
 180 with  $100 \times 10^3$  neurons running for  $10 \times 10^3$  simulation timesteps, required less than 120 MB – a small  
 181 fraction of the memory on a modern GPU. While efficiently handling spikes stored in a bitfield is a little  
 182 trickier than working with a list of neuron indices, GeNN provides an efficient C++ helper function for  
 183 saving the spikes stored in a bitfield to a text file and a numpy-based method for decoding them in PyGeNN.



**Figure 3.** Simulation times of the microcircuit model running on various GPU hardware for 1 s of biological time. ‘Overhead’ refers to time spent in simulation loop but not within CUDA kernels. The dotted horizontal line indicates realtime performance



## 184 2.5 Cortical microcircuit model

185 Potjans and Diesmann (2014) developed this model of  $1\text{ mm}^3$  of early-sensory cortex. The model  
 186 consists of 77 169 LIF neurons, divided into separate populations representing the excitatory and inhibitory  
 187 population in each of 4 cortical layers (2/3, 4, 5 and 6) as illustrated by figure 1. Neurons in each population  
 188 are connected randomly with numbers of synapses derived from an extensive review of the anatomical  
 189 literature. Within each synaptic projection, all synaptic strengths and transmission delays are normally  
 190 distributed and, in total, the model has approximately  $0.3 \times 10^9$  synapses. As well as receiving synaptic  
 191 input, each neuron in the network also receives an independent Poisson input current, representing input  
 192 from neighbouring un-modelled cortical regions. For a full description of the model parameters, please  
 193 refer to Potjans and Diesmann (2014, tables 4 and 5) and for a description of the strategies used by GeNN to  
 194 parallelise the initialisation and subsequent simulation of this network, please refer to Knight and Nowotny  
 195 (2018, section 2.3). This model requires simulation using a relatively small timestep of 0.1 ms, making the  
 196 overheads of copying spikes from the GPU every timestep particularly significant.

## 197 2.6 Random balanced network HPC benchmark model

198 sdasd

## 3 RESULTS

199 In the following subsection we will analyse the performance of the models introduced in sections 2.5 and 2.6  
 200 on a representative selection of NVIDIA GPU hardware:

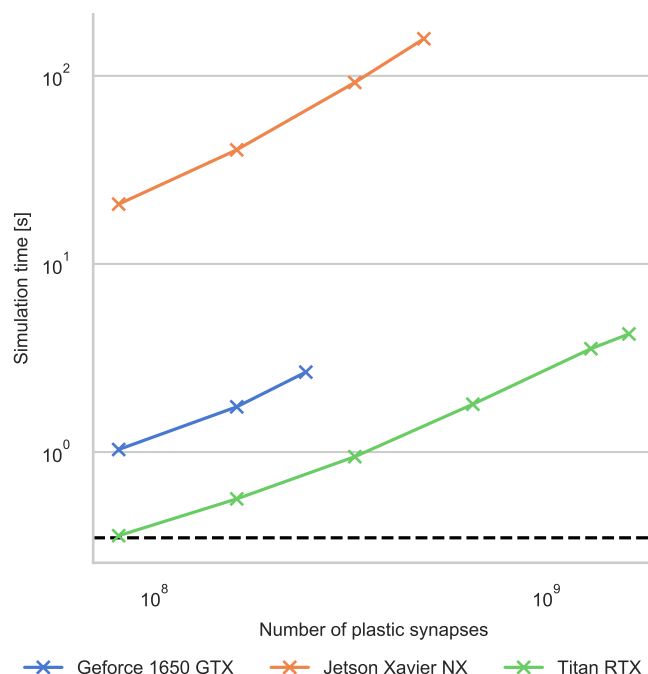
- 201 • Jetson Xavier NX – a low-power embedded system with 8 GB of shared memory.
- 202 • GeForce GTX 1650 – a low-end desktop GPU with 4 GB of dedicated memory.
- 203 • Titan RTX – a high-end workstation GPU with 24 GB of dedicated memory.

### 204 3.1 Cortical microcircuit model performance

205 Figure 3 shows the simulation times for the full-scale microcircuit mode and, as one might predict, the  
 206 Jetson Xavier NX is slower than the two desktop GPUs. However, considering that it only consumes a  
 207 maximum of 15 W compared to 75 W or 320 W for the GeForce GTX 1650 and Titan RTX respectively, it  
 208 still performs impressively.

209 The time taken to actually simulate the models ('Neuron simulation' and 'Synapse simulation') are the  
 210 same when using Python and C++ as all GeNN optimisation options are exposed to PyGeNN. However,  
 211 both the PyGeNN and C++ simulations spend a significant amount of every simulation step copying spike  
 212 data off the device and storing it in a suitable data structure ('Overhead'). Because Python is an interpreted  
 213 language, such operations are inherently slower – this is particularly noticeable on devices with a slower  
 214 CPU such as the Jetson Xavier NX. Unlike the desktop GPUs, the Jetson Xavier NX's 8 GB of memory is  
 215 shared between the GPU and the CPU meaning that data doesn't have to be copied between GPU and CPU  
 216 memory and can instead be accessed by both. **(TODO: RUN XAVIER NX WITHOUT RECORDING AND**  
 217 **WITHOUT ZERO COPY)** While, using this shared memory for recording spikes, reduces the overheads  
 218 associated with copying data off the device, because the GPU and CPU caches are not coherent, caching  
 219 must be disabled on this memory which reduces the performance of the neuron kernel. However, when  
 220 the spike recording system described in section 2.4 is used, spike data is kept in GPU memory until the

end of the simulation and this overhead is reduced by around a factor of 10. This now means that, on the high-end desktop GPU, this simulation now runs faster than real-time – previously only achievable using a specialised neuromorphic system (Rhodes et al., 2020) and significantly faster than other recently published GPU simulators (Golosio et al., 2020).



**Figure 4.**

## 4 DISCUSSION

discuss!

- PyGeNN as an intermediate layer - PyNN, ML
- Cost of C++ - Python calls in models

## CONFLICT OF INTEREST STATEMENT

The authors declare that the research was conducted in the absence of any commercial or financial relationships that could be construed as a potential conflict of interest.

## AUTHOR CONTRIBUTIONS

JK and TN wrote the paper. TN is the original developer of GeNN. AK was the original developer of PyGeNN. JK is currently the primary developer of both GeNN and PyGeNN and was responsible for implementing the spike recording system. JK performed the experiments and the analysis of the results that are presented in this work.



## FUNDING

234 This work was funded by the EPSRC (Brains on Board project, grant number EP/P006094/1).

## ACKNOWLEDGMENTS

235 This is a short text to acknowledge the contributions of specific colleagues, institutions, or agencies that  
236 aided the efforts of the authors.

## DATA AVAILABILITY STATEMENT

237 The datasets [GENERATED/ANALYZED] for this study can be found in the [NAME OF REPOSITORY]  
238 [LINK].

## REFERENCES

- 239 Akar, N. A., Cumming, B., Karakasis, V., Kusters, A., Klijn, W., Peyser, A., et al. (2019). Arbor — A  
240 Morphologically-Detailed Neural Network Simulation Library for Contemporary High-Performance  
241 Computing Architectures. In *2019 27th Euromicro International Conference on Parallel, Distributed  
242 and Network-Based Processing (PDP)* (IEEE), 274–282. doi:10.1109/EMPDP.2019.8671560
- 243 Balaji, A., Adiraju, P., Kashyap, H. J., Das, A., Krichmar, J. L., Dutt, N. D., et al. (2020). PyCARL: A  
244 PyNN Interface for Hardware-Software Co-Simulation of Spiking Neural Network
- 245 Beazley, D. M. (1996). Using SWIG to control, prototype, and debug C programs with Python. In *Proc.  
246 4th Int. Python Conf*
- 247 Buzsáki, G. and Mizuseki, K. (2014). The log-dynamic brain: how skewed distributions affect network  
248 operations. *Nature reviews. Neuroscience* 15, 264–78. doi:10.1038/nrn3687
- 249 Carnevale, N. T. and Hines, M. L. (2006). *The NEURON book* (Cambridge University Press)
- 250 Chou, T.-s., Kashyap, H. J., Xing, J., Listopad, S., Rounds, E. L., Beyeler, M., et al. (2018). CARLsim 4:  
251 An Open Source Library for Large Scale, Biologically Detailed Spiking Neural Network Simulation  
252 using Heterogeneous Clusters. In *2018 International Joint Conference on Neural Networks (IJCNN)*  
253 (IEEE), 1–8. doi:10.1109/IJCNN.2018.8489326
- 254 Davison, A. P., Brüderle, D., Eppler, J., Kremkow, J., Müller, E., Pecevski, D., et al. (2008). PyNN: A  
255 Common Interface for Neuronal Network Simulators. *Frontiers in neuroinformatics* 2, 11. doi:10.3389/  
256 neuro.11.011.2008
- 257 Eppler, J. M., Helias, M., Müller, E., Diesmann, M., and Gewaltig, M. O. (2009). PyNEST: A convenient  
258 interface to the NEST simulator. *Frontiers in Neuroinformatics* 2, 1–12. doi:10.3389/neuro.11.012.2008
- 259 Gewaltig, M.-O. and Diesmann, M. (2007). NEST (NEural Simulation Tool). *Scholarpedia* 2, 1430
- 260 Golosio, B., Tiddia, G., De Luca, C., Pastorelli, E., Simula, F., and Paolucci, P. S. (2020). A new GPU  
261 library for fast simulation of large-scale networks of spiking neurons , 1–27
- 262 Hines, M. L., Davison, A. P., and Müller, E. (2009). NEURON and Python. *Frontiers in Neuroinformatics*  
263 3, 1–12. doi:10.3389/neuro.11.001.2009
- 264 Hunter, J. D. (2007). Matplotlib: A 2D graphics environment. *Computing in Science & Engineering* 9,  
265 90–95. doi:10.1109/MCSE.2007.55
- 266 Knight, J. C. and Nowotny, T. (2018). GPUs Outperform Current HPC and Neuromorphic Solutions  
267 in Terms of Speed and Energy When Simulating a Highly-Connected Cortical Model. *Frontiers in  
268 Neuroscience* 12, 1–19. doi:10.3389/fnins.2018.00941

- 269 Knight, J. C. and Nowotny, T. (2020). Larger GPU-accelerated brain simulations with procedural  
270 connectivity. *bioRxiv* doi:10.1101/2020.04.27.063693
- 271 Millman, K. J. and Aivazis, M. (2011). Python for scientists and engineers. *Computing in Science and*  
272 *Engineering* 13, 9–12. doi:10.1109/MCSE.2011.36
- 273 Potjans, T. C. and Diesmann, M. (2014). The Cell-Type Specific Cortical Microcircuit: Relating Structure  
274 and Activity in a Full-Scale Spiking Network Model. *Cerebral Cortex* 24, 785–806. doi:10.1093/cercor/  
275 bhs358
- 276 Rhodes, O., Peres, L., Rowley, A. G. D., Gait, A., Plana, L. A., Brenninkmeijer, C., et al. (2020). Real-time  
277 cortical simulation on neuromorphic hardware. *Philosophical Transactions of the Royal Society A:*  
278 *Mathematical, Physical and Engineering Sciences* 378, 20190160. doi:10.1098/rsta.2019.0160
- 279 Stimberg, M., Brette, R., and Goodman, D. F. (2019). Brian 2, an intuitive and efficient neural simulator.  
280 *eLife* 8, 1–41. doi:10.7554/eLife.47314
- 281 Stimberg, M., Goodman, D. F., and Nowotny, T. (2020). Brian2GeNN: accelerating spiking neural network  
282 simulations with graphics hardware. *Scientific Reports* 10, 1–12. doi:10.1038/s41598-019-54957-7
- 283 Van Der Walt, S., Colbert, S. C., and Varoquaux, G. (2011). The NumPy array: A structure for efficient  
284 numerical computation. *Computing in Science and Engineering* 13, 22–30. doi:10.1109/MCSE.2011.37
- 285 Yavuz, E., Turner, J., and Nowotny, T. (2016). GeNN: a code generation framework for accelerated brain  
286 simulations. *Scientific reports* 6, 18854. doi:10.1038/srep18854